

# Cryptography

Stream Ciphers


# Content



Content
Introduction
Hardware-Oriented Stream Ciphers
Software-Oriented Stream Ciphers


# Introduction

- Symmetric ciphers:
  - **Block ciphers**: mix blocks of plaintext bits to produce blocks of ciphertext bits.
  - **Stream ciphers**: generate random bits from the key and XOR it with the plaintext.
    - Similar to the OTP.
    - Akin to DRBG, not PRNG.



# Introduction

- Stream ciphers take two inputs:
  - **Key**: secret – its size 128 or 256 bits.
  - **Nonce**: Not secret – its size is between 64 and 128 bits




- Stream ciphers produce a pseudorandom stream of bits called the **keystream**.
- **Encryption**: XOR the plaintext with the keystream.
- **Decryption**: XOR the ciphertext with the keystream

# Introduction

- Never use the same key and nonce pair twice.
  - Nonce = number used once
- If you have  $(K_1, N_1)$  encrypting a message  $P_1$ .
- You can use another pair  $(K_1, N_2)$ ,  $(K_2, N_1)$  or  $(K_2, N_2)$  to encrypt another message  $P_2$ .
- Using the same key and nonce will produce the same keystream.
  - If you have  $C_1 = P_1 \oplus KS$  and  $C_2 = P_2 \oplus KS$  and you know  $P_1$ , you can get  $P_2 = C_1 \oplus C_2 \oplus P_1$

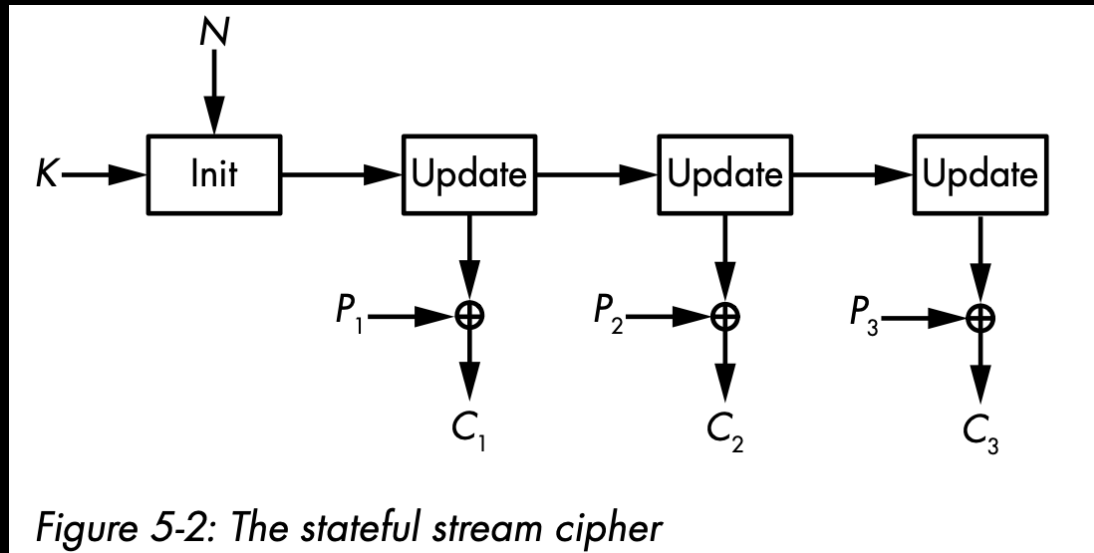
# Introduction



# Introduction

## Stateful stream ciphers

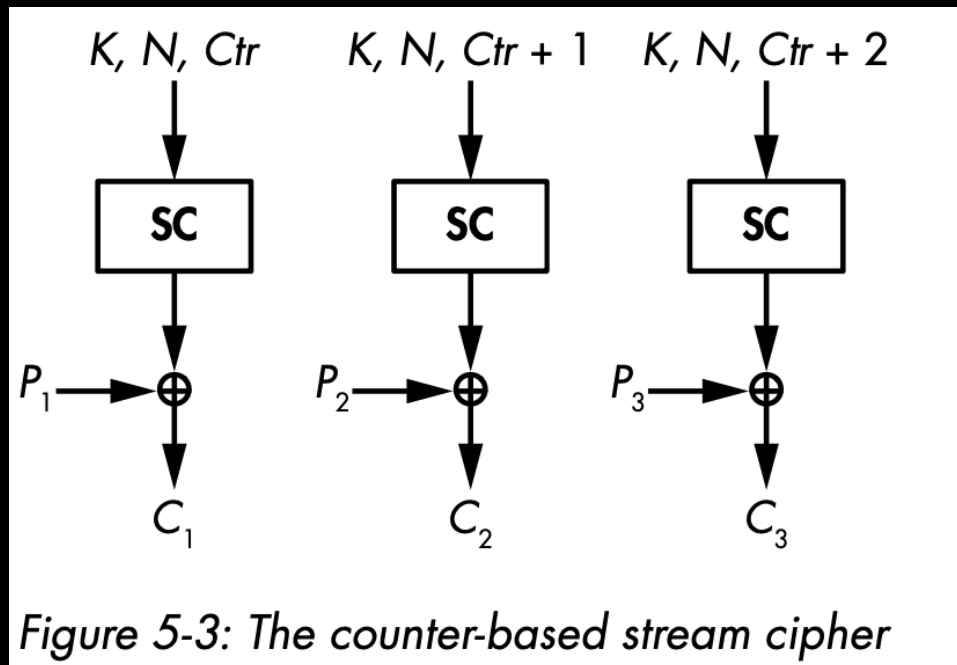
- Have a secret internal state that evolves throughout keystream generation.
  1. The cipher initializes the state from the key and the nonce.
  2. Calls an *update* function to update the state value.
  3. Produce keystream bits from the state
- Example: **RC4**



# Introduction

## Counter-based stream ciphers



- Inputs a key, a nonce, and a counter and produce chunks of keystream.
- Example: **Salsa20**





# Introduction

- Stream ciphers can be either:
  - Hardware-oriented stream ciphers.
    - Application Specific Integrated Circuit (ASIC)
    - Programmable Logic Device (PLD)
    - Field Programmable Gate Arrays (FPGA)
  - Software-oriented stream ciphers.



# Content

## Content

Introduction

Hardware-Oriented Stream Ciphers

Software-Oriented Stream Ciphers

# Hardware-Oriented Stream Ciphers


- Electronic circuit that implements a cryptographic algorithm.
  - Lower cost on hardware vs block ciphers
  - Less memory, smaller chip
- Hardware stream ciphers rely on **Feedback Shift Registers (FSR)**.

# Hardware-Oriented Stream Ciphers

- Components of feedback shift register (FSR):
  - **State  $R$** : an array of bits, or registers.
  - **Feedback function  $f$** : updates the state and output bits
- FSR works as follows:
  - Shift the current state to the left by 1-bit
  - Return the shifted bit as output
  - Fill the empty position with the bit  $f(R)$

$$R_{t+1} = (R_t \ll 1) \text{ OR } f(R_t)$$

# Hardware-Oriented Stream Ciphers



# Hardware-Oriented Stream Ciphers

- Example: consider a 4-bit FSR whose feedback function  $f$  XORs all 4 bits together. The initial state is

1	1	0	0
---	---	---	---

- Shift the bits to the left and output the shifted bit

Output: 1

1	0	0	x
---	---	---	---

- Fill the rightmost position with the bit  $f(1100) = 1 \oplus 1 \oplus 0 \oplus 0 = 0$

1	0	0	0
---	---	---	---

# Hardware-Oriented Stream Ciphers

- The next update is as follows

Output: 11

1	0	0	0
---	---	---	---

0	0	0	x
---	---	---	---

- Fill the rightmost position with the bit  $f(1000) = 1 \oplus 0 \oplus 0 \oplus 0 = 1$

0	0	0	1
---	---	---	---

# Hardware-Oriented Stream Ciphers

- The next 3 updates return three 0 bits and give the following state values:

Output: 110      

0	0	1	1
---	---	---	---

Output: 1100      

0	1	1	0
---	---	---	---

Output: 11000      

1	1	0	0
---	---	---	---



# Hardware-Oriented Stream Ciphers

- The next 3 updates return three 0 bits and give the following state values:

Output: 110      

0	0	1	1
---	---	---	---

Output: 1100      

0	1	1	0
---	---	---	---

Output: 11000      




1	1	0	0
---	---	---	---

- We return to the initial state after 5 iterations!
- The output is the same as the initial state!

# Hardware-Oriented Stream Ciphers




- We say that 5 is the *period* of the FSR.
  - FSR period: the number of updates needed until the FSR enters the same state again.
- So, if we clock the register 20 times, the output will be:  
11000110001100011000
- Patterns are bad, avoid them.
- Short periods → predictable patterns → bad FSRs

# Hardware-Oriented Stream Ciphers




What is the period of 0000?

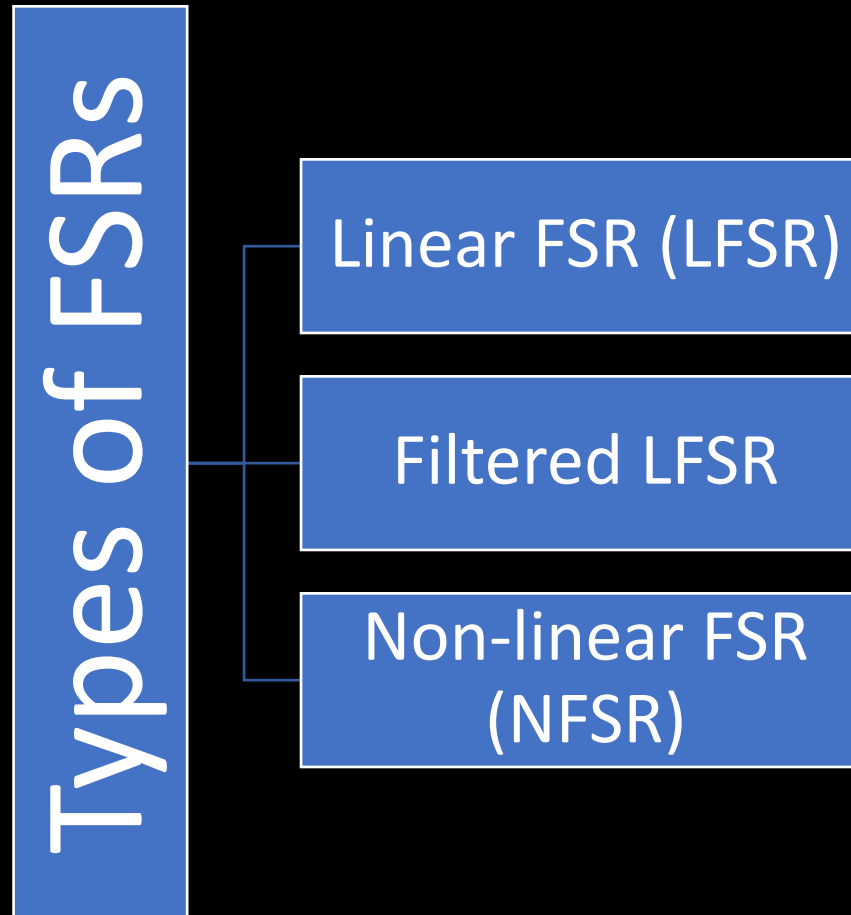
# Hardware-Oriented Stream Ciphers



What is the period of 0000?



# Hardware-Oriented Stream Ciphers

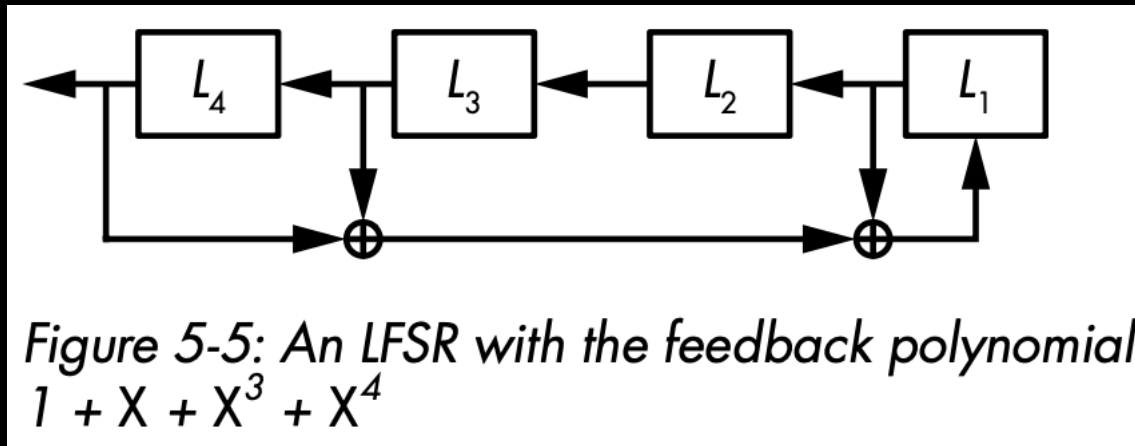


# Hardware-Oriented Stream Ciphers

- LFSRs apply linear feedback function.
  - E.g., a function that do XOR of the bits of the state.
- In crypto, linear  $\rightarrow$  patterns  $\rightarrow$  predictable  $\rightarrow$  bad.
- The choice of which bits are XORed together is crucial.
  - We want to guarantee a maximal period.
  - For an  $n$ -bit register, the maximal period is  $2^n - 1$ .

# Hardware-Oriented Stream Ciphers

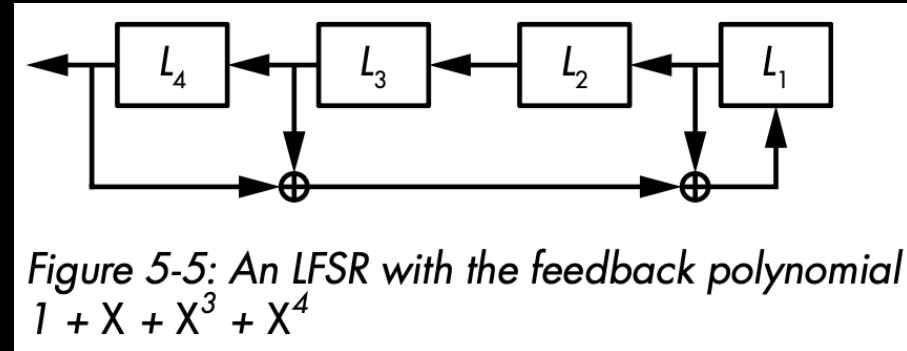
- LFSRs are represented as polynomials:  $1 + X + X^2 + \dots + X^n$ .
  - The term  $X^i$  is the  $i$ th bit that is XORed in the feedback function.
- Example: the 4-bit LFSR is represented as  $1 + X + X^3 + X^4$  because the bits 1, 3, and 4 are XORed in the feedback function.



# Hardware-Oriented Stream Ciphers

- The period of this LFSR is not maximal.
  - We get the initial state after six updates

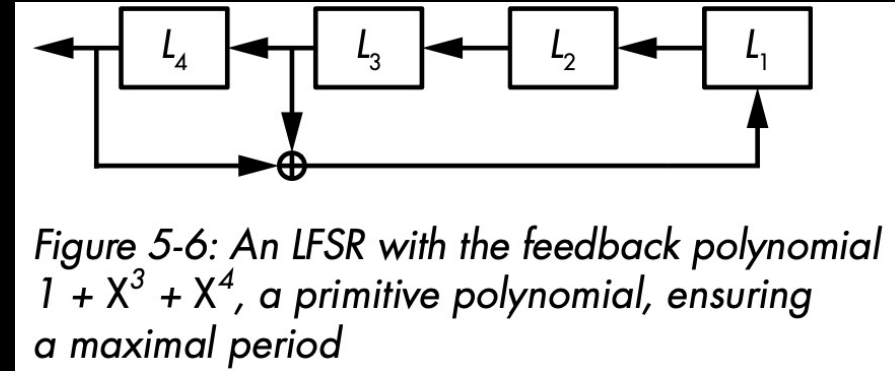
Initial state → 0 0 0 1  
0 0 1 1  
0 1 1 1  
1 1 1 0  
1 1 0 0  
1 0 0 0  
0 0 0 1





# Hardware-Oriented Stream Ciphers


- This LFSR has maximal period
  - $1 + X^3 + X^4$
  - The initial value: 0 0 0 1



0 0 0 1	0 0 1 1	0 1 0 1	1 1 1 0
0 0 1 0	0 1 1 0	1 0 1 1	1 1 0 0
0 1 0 0	1 1 0 1	0 1 1 1	1 0 0 0
1 0 0 1	1 0 1 0	1 1 1 1	0 0 0 1

# Hardware-Oriented Stream Ciphers

- Filtered LFSR: hides the linearity of the LFSR by using nonlinear function.
  - The output of the LFSR is passed to a nonlinear function before producing the output



- $g$  is a nonlinear function that applies XORs with ANDs or ORs

# Hardware-Oriented Stream Ciphers

- Complex attacks can break the filtered LFSR
  - *Patching a broken algorithm with a slightly stronger layer won't make the whole thing secure.*
- Attacks:
  - **Algebraic attacks:** solve the nonlinear equations deduced from the output bits.
  - **Cube attacks:** compute derivatives of the nonlinear equations to reduce the degree of the system down to one and then solve it.
  - **Fast correlation attacks:** exploit filtering functions that, despite their nonlinearity, tend to behave like linear functions.

# Hardware-Oriented Stream Ciphers

- Nonlinear FSR: like LFSRs but with a nonlinear feedback function.
  - Uses ANDs, ORs and XORs


- Example: a 4-bit NFSR feedback function

$$N_1 + N_2 + N_1N_2 + N_3N_4$$

- **Downside:** cannot determine if its period is maximal
  - For  $n$ -bit NFSR, we need to run  $2^n$  trials; impossible for large  $n$  (e.g.,  $n = 80$ )

# Hardware-Oriented Stream Ciphers

- Grain-128a: a stream cipher included in the eSTREAM 2008 competition.



- Combines 128-bit LFSR, 128-bit NFSR, and a filter function  $h$
- The LFSR has a maximal period of at least  $2^{128} - 1$

# Hardware-Oriented Stream Ciphers

How Grain-128a works:

**1. Input:** 128-bit key and 96-bit nonce

1. The 128-bit key goes to the NFSR, the 96-bit nonce goes to the LFSR
2. The LFSR is padded with 31 ones and 1 zero bits

**2. Initialization phase:** updates the system 256 times before returning the first keystream bit.

**3. LFSR feedback function:**

$$\mathbf{f}(L) = L_{32} + L_{47} + L_{58} + L_{90} + L_{121} + L_{128}$$

**4. NFSR feedback function:**

$$\begin{aligned} \mathbf{g}(N) = & N_{32} + N_{37} + N_{72} + N_{102} + N_{128} + N_{44}N_{60} + N_{61}N_{125} + N_{63}N_{67} + N_{69}N_{101} \\ & + N_{80}N_{88} + N_{110}N_{111} + N_{115}N_{117} + N_{46}N_{50}N_{58} + N_{103}N_{104}N_{106} + N_{33}N_{35}N_{36}N_{40} \end{aligned}$$

# Hardware-Oriented Stream Ciphers

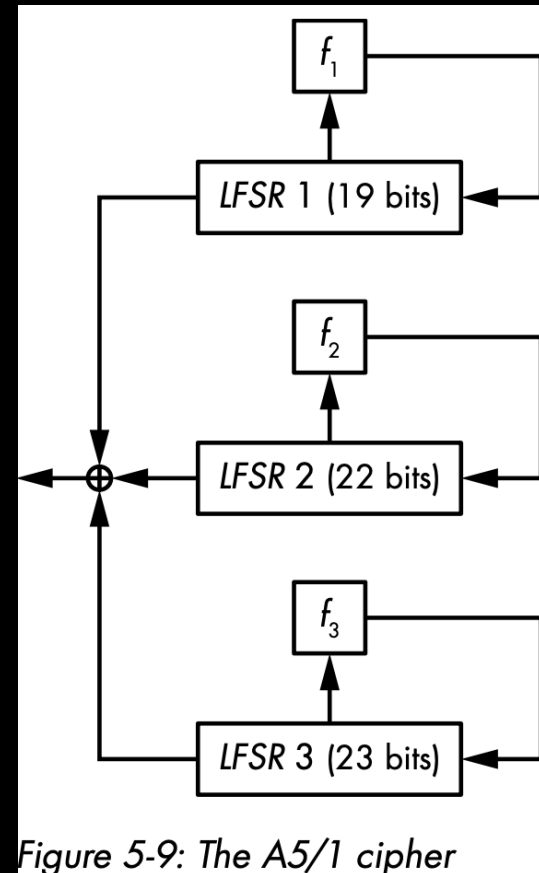
How Grain-128a works:

7. **Nonlinear  $h$** : takes 9 bits from the NFSR and 7 bits from the LFSR and combines them.

- Used in low-level embedded systems:
  - Fast and secure
  - Proprietary


# Hardware-Oriented Stream Ciphers

- **A5/1 cipher**: encrypts voice communications in the 2G mobile standard.
  - Broken
- It uses three LFSRs
- Key: 64 bits
- Nonce: 22 bits
  - Changed every time frame





# Hardware-Oriented Stream Ciphers



# Hardware-Oriented Stream Ciphers

- **Subtle attacks:**

- Called a *guess-and-determine* attack
- An attacker guesses certain secret values of the state to determine others

```
For all  $2^{19}$  values of LFSR 1's initial state
  For all  $2^{22}$  values of LFSR 2's initial state
    For all  $2^{11}$  values of LFSR 3's clocking bit during the first 11 clocks
      Reconstruct LFSR 3's initial state
      Test whether guess is correct; if yes, return; else continue
```

- This includes solving equations that depends on the bits guessed
- Takes at most  $2^{19} \times 2^{22} \times 2^{11} = 2^{52}$  operations instead of  $2^{64}$

# Hardware-Oriented Stream Ciphers

- **Brutal attacks:**

- Time-memory-trade-off (TMTO) attack
- **Codebook attack** – precompute a table of  $2^{64}$  of key-value pairs and store them
- The attack is fast – just look up a value in memory
- Takes very long to compute the table, requires large memory, 256 exabytes
- **TMTO attacks reduce the memory required by a codebook attack** at the price of increased computation during the online phase of the attack
- In 2010, researchers took about 2 months to generate two terabytes' tables, using GPUs and running 100,000 instances of A5/1 in parallel.

# Hardware-Oriented Stream Ciphers

## TASKS

1. Implement the LFSR class that takes an input an initial state and the bits position to XOR. Try with seed = [1, 0, 1, 1] and positions = [0, 2].
2. Implement a function that computes the period of the previous LFSR and its feedback function. What is the period of each of these feedback functions [0, 2, 3] and [0, 3] on the state [1, 0, 0, 0]?
3. Implement the NFSR class that takes an input an initial state and the bits position. The feedback bit is computed using:  $N_2 + N_3$  and the output bit is computed using:  $N_0 + N_1 + N_0N_1 + N_2N_3$

# Content

## Content

Introduction

Hardware-Oriented Stream Ciphers

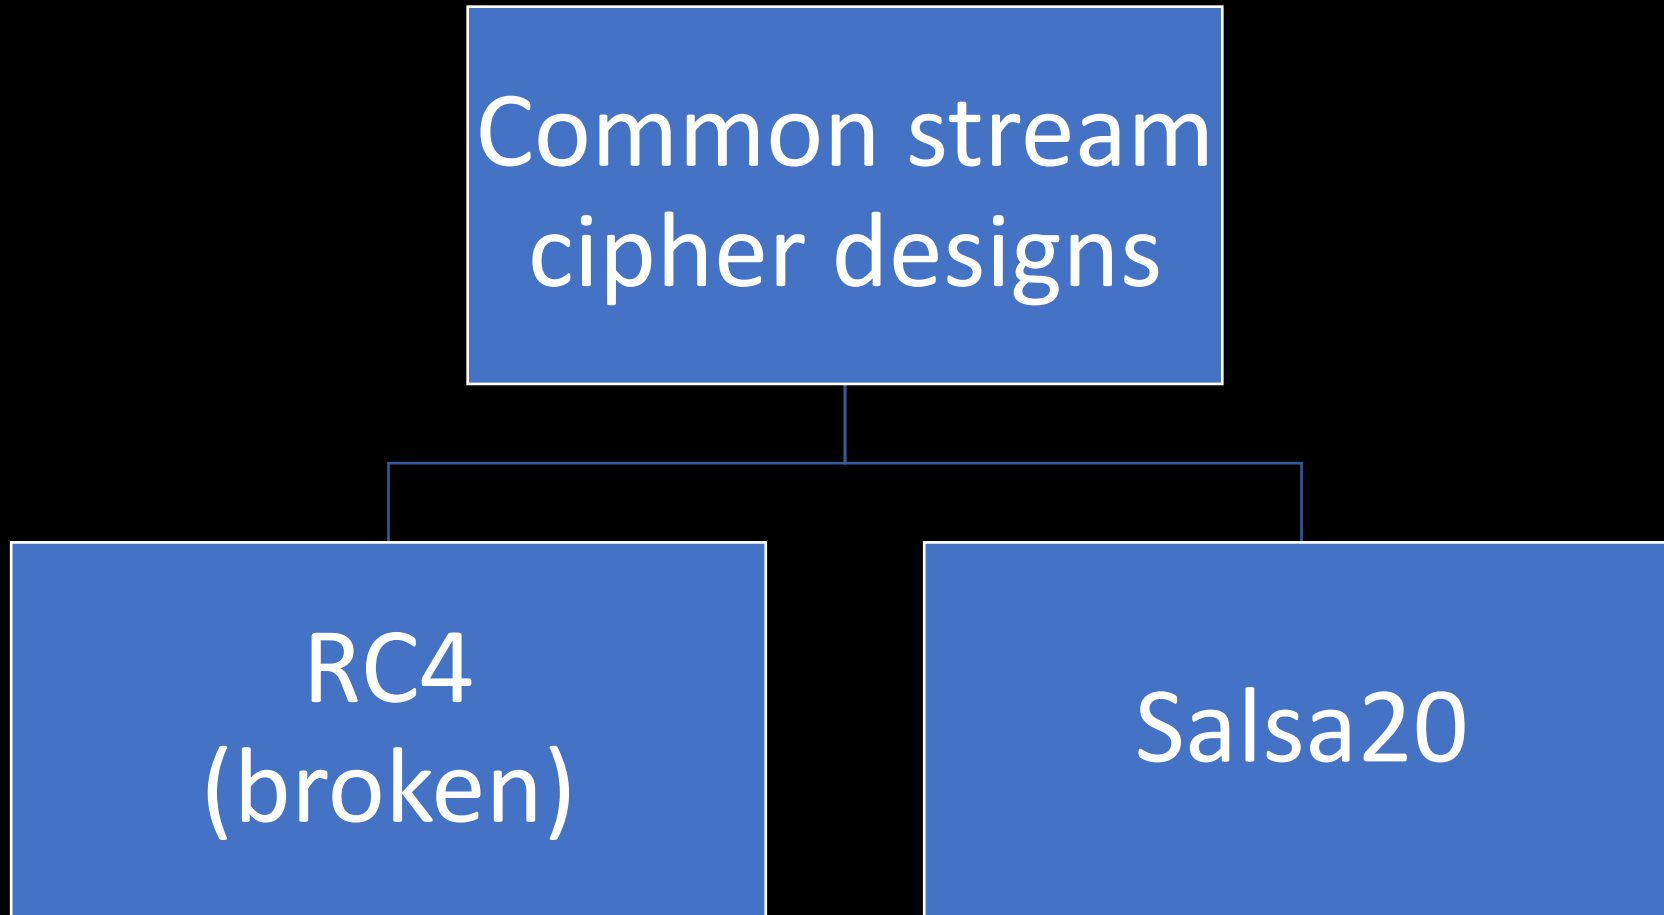


Software-Oriented Stream Ciphers

# Software-Oriented Stream Ciphers

- Work with bytes, 32-bit, or 64-bit words instead of individual bits.
- Better than hardware ciphers when run as part of a software.
  - Browsers and servers
- Popularity of software stream ciphers:
  - Some 4G networks use stream ciphers that work with 32-bit words.
    - SONW3G and ZUC
  - Some block ciphers are weak against attacks
    - e.g., padding oracle attack against block ciphers in CBC mode
  - Easier to specify and to implement than block cipher
    - AES (block cipher) + CTR mode of operation → stream cipher

# Software-Oriented Stream Ciphers



# Software-Oriented Stream Ciphers

## RC4

- RC4 was used in many applications:
  - WiFi encryption standard (WEP)
  - Transport Layer Security (TLS) protocol for establishing HTTPS connections.
- How RC4 works: Just swaps bytes (no complex operations)
  - It has an internal array,  $S$ , of size 256 bytes:  $S[0] = 0, S[1] = 1, \dots, S[255] = 255$
  - $S$  is initialized from a key,  $k$ , according to a key scheduling algorithm (KSA)

```
j = 0
# set S to the array S[0] = 0, S[1] = 1, . . . , S[255] = 255
S = range(256)
# iterate over i from 0 to 255
for i in range(256):
    # compute the sum of v
    j = (j + S[i] + K[i % n]) % 256
    # swap S[i] and S[j]
    S[i], S[j] = S[j], S[i]
```



# Software-Oriented Stream Ciphers

## RC4

- The internal state  $S$  looks like this before executing the KSA:  
0, 1, 2, 3, ..., 255
- After executing the KSA, it will become  
0, 35, 90, 116, ..., 4
- If you change one bit of the key, you will get a totally different state  
32, 11, 10, 212, ..., 2

# Software-Oriented Stream Ciphers

## RC4

- After initializing the state  $S$ , RC4 generates the key stream  $KS$ .
  - $\text{Size}(KS) = \text{size}(\text{plaintext})$
- The ciphertext is computed as  $c = KS \oplus \text{plaintext}$
- The keystream is computed as
  - $m$  is the size of the plaintext in bytes

```
i = 0
j = 0
for b in range(m):
    i = (i + 1) % 256
    j = (j + S[i]) % 256
    S[i], S[j] = S[j], S[i]
    KS[b] = S[(S[i] + S[j]) % 256]
```

# Software-Oriented Stream Ciphers

## RC4

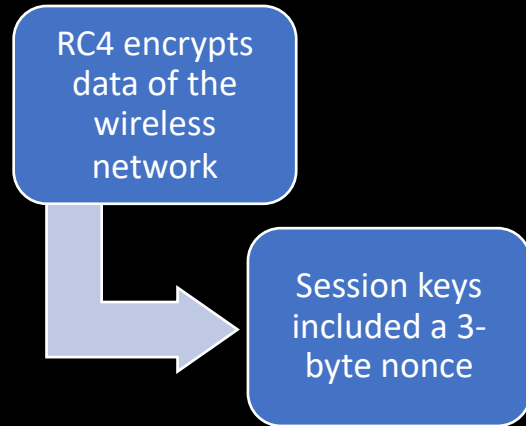
- RC4 in WEP

RC4 encrypts  
data of the  
wireless  
network

# Software-Oriented Stream Ciphers

## RC4

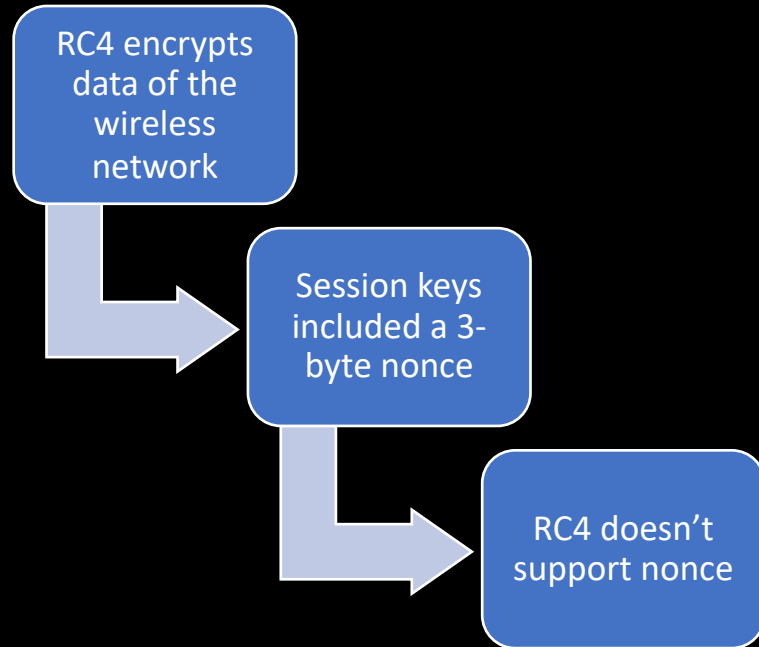
- RC4 in WEP



# Software-Oriented Stream Ciphers

## RC4

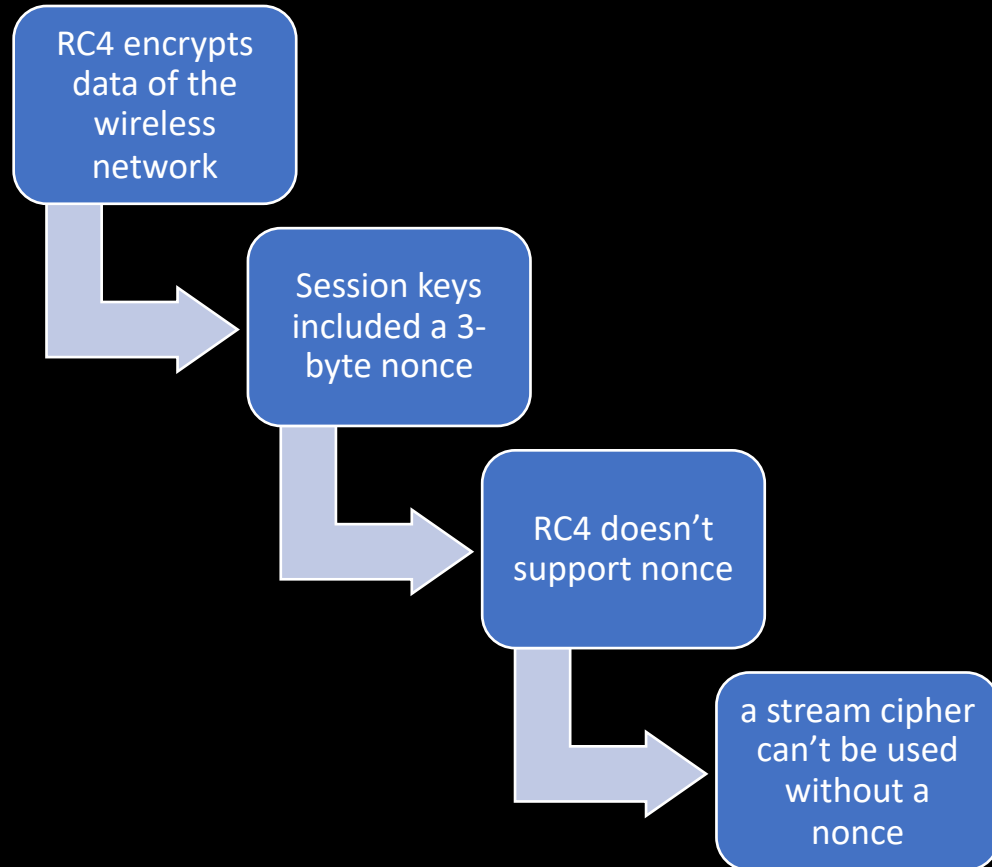
- RC4 in WEP



# Software-Oriented Stream Ciphers

## RC4

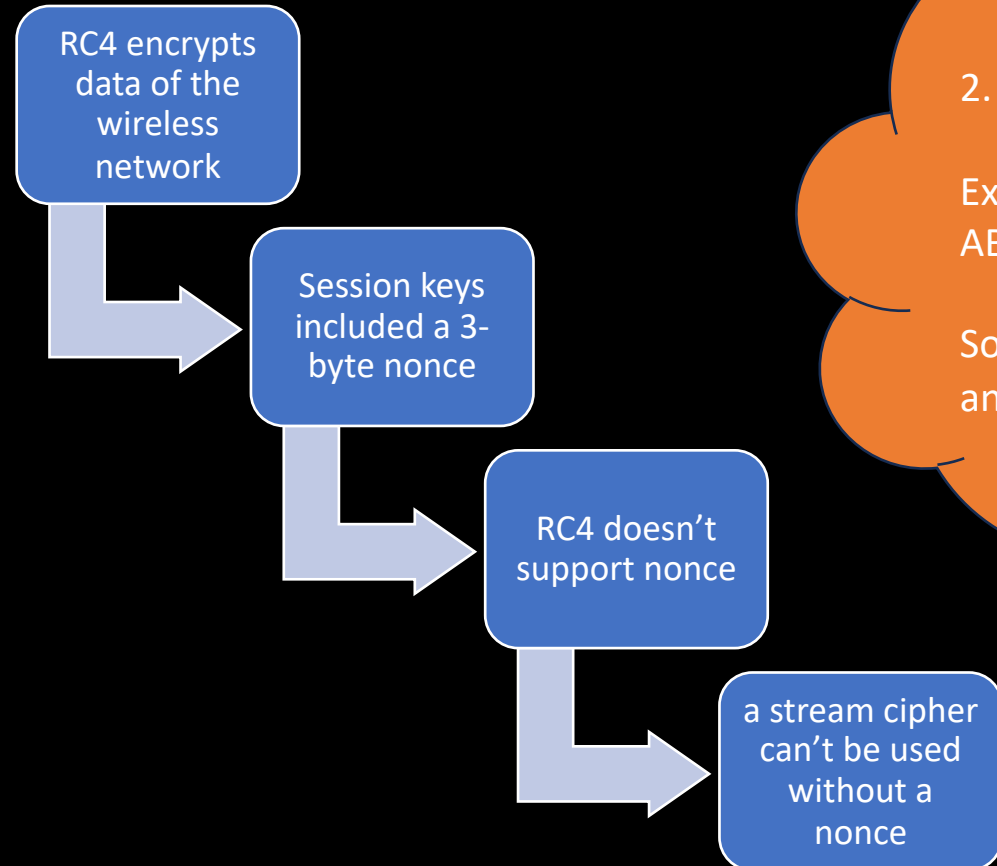
- RC4 in WEP



# Software-Oriented Stream Ciphers

## RC4

- RC4 in WEP



1. Include a 24-bit nonce in the wireless frame's header
2. Prepend it to the WEP key.

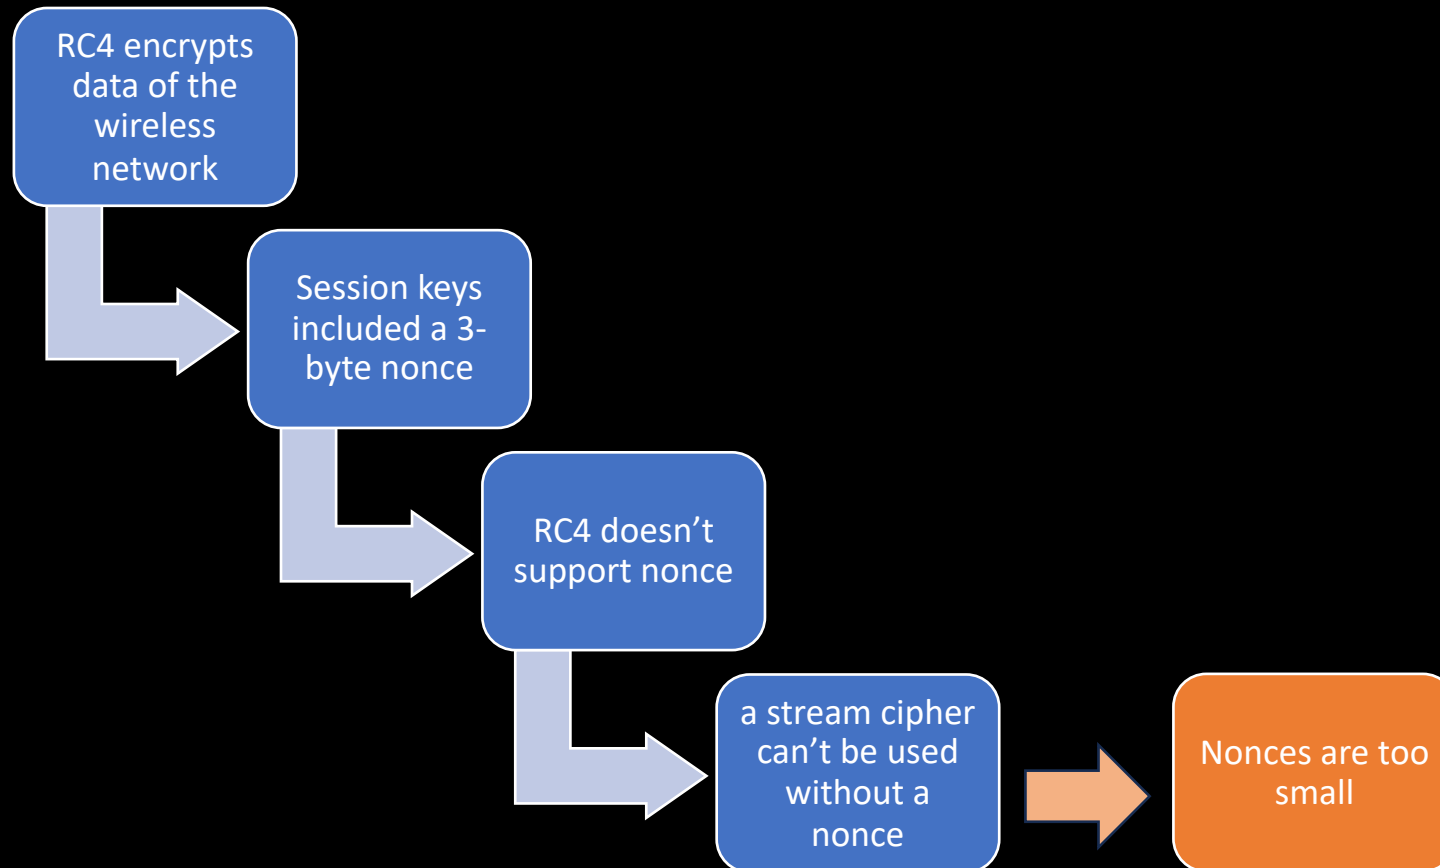
Ex: if the nonce is 123 and the key is ABC, then the RC4 key is 123ABC.

So, 40-bit keys, yield 64-bit security, and 104-bit keys, yield 128-bit key

# Software-Oriented Stream Ciphers

## RC4

- RC4 in WEP

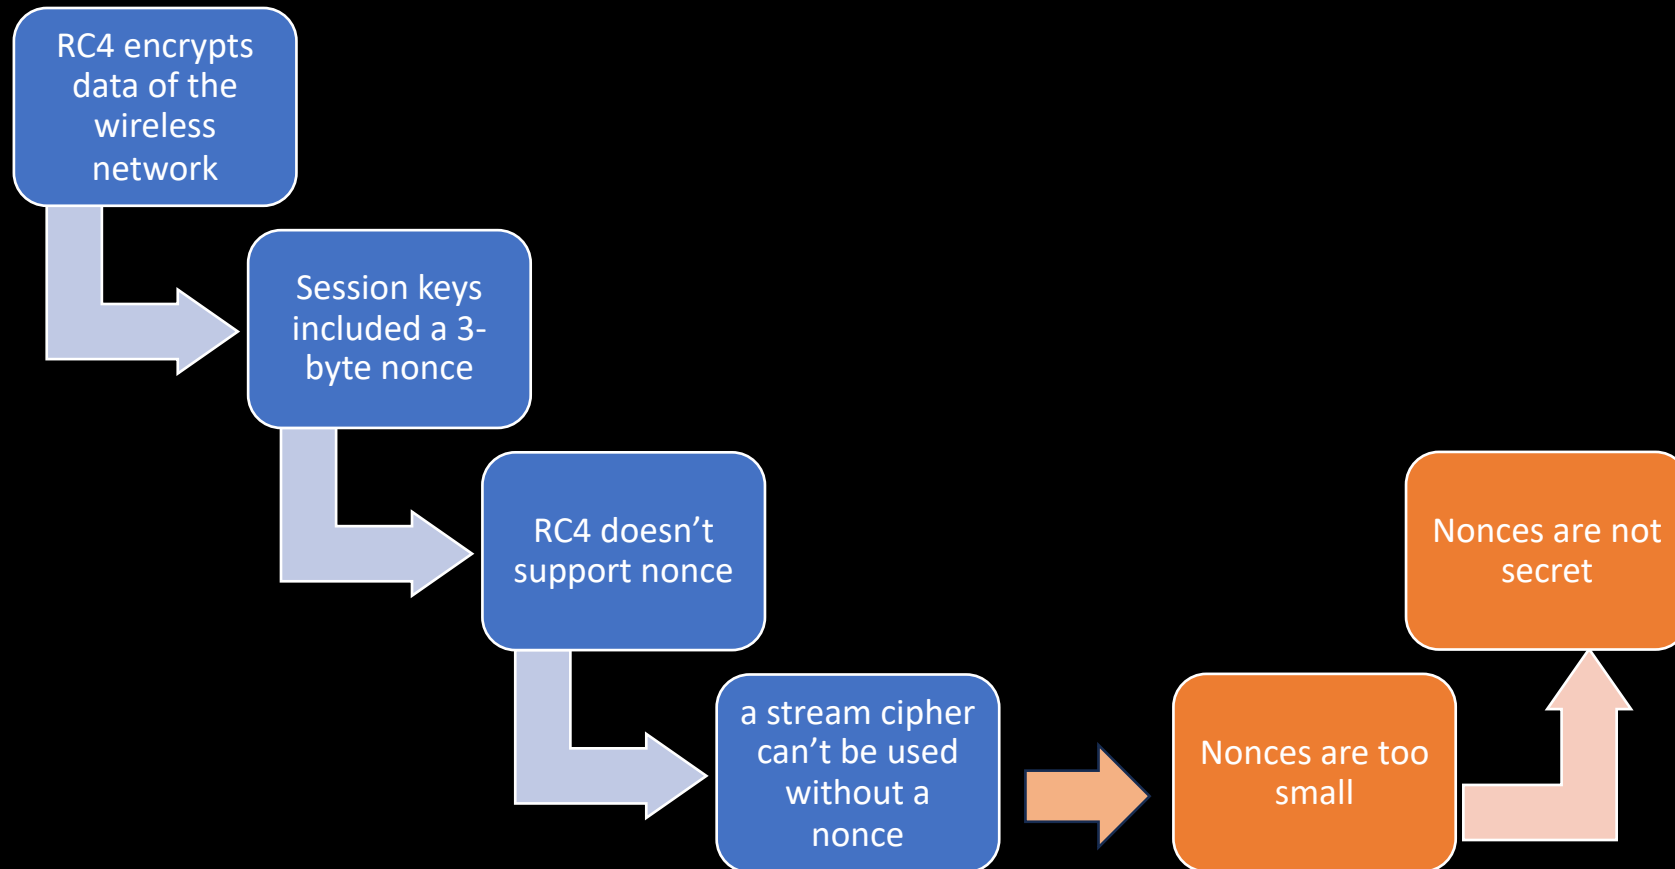




# Software-Oriented Stream Ciphers

## RC4

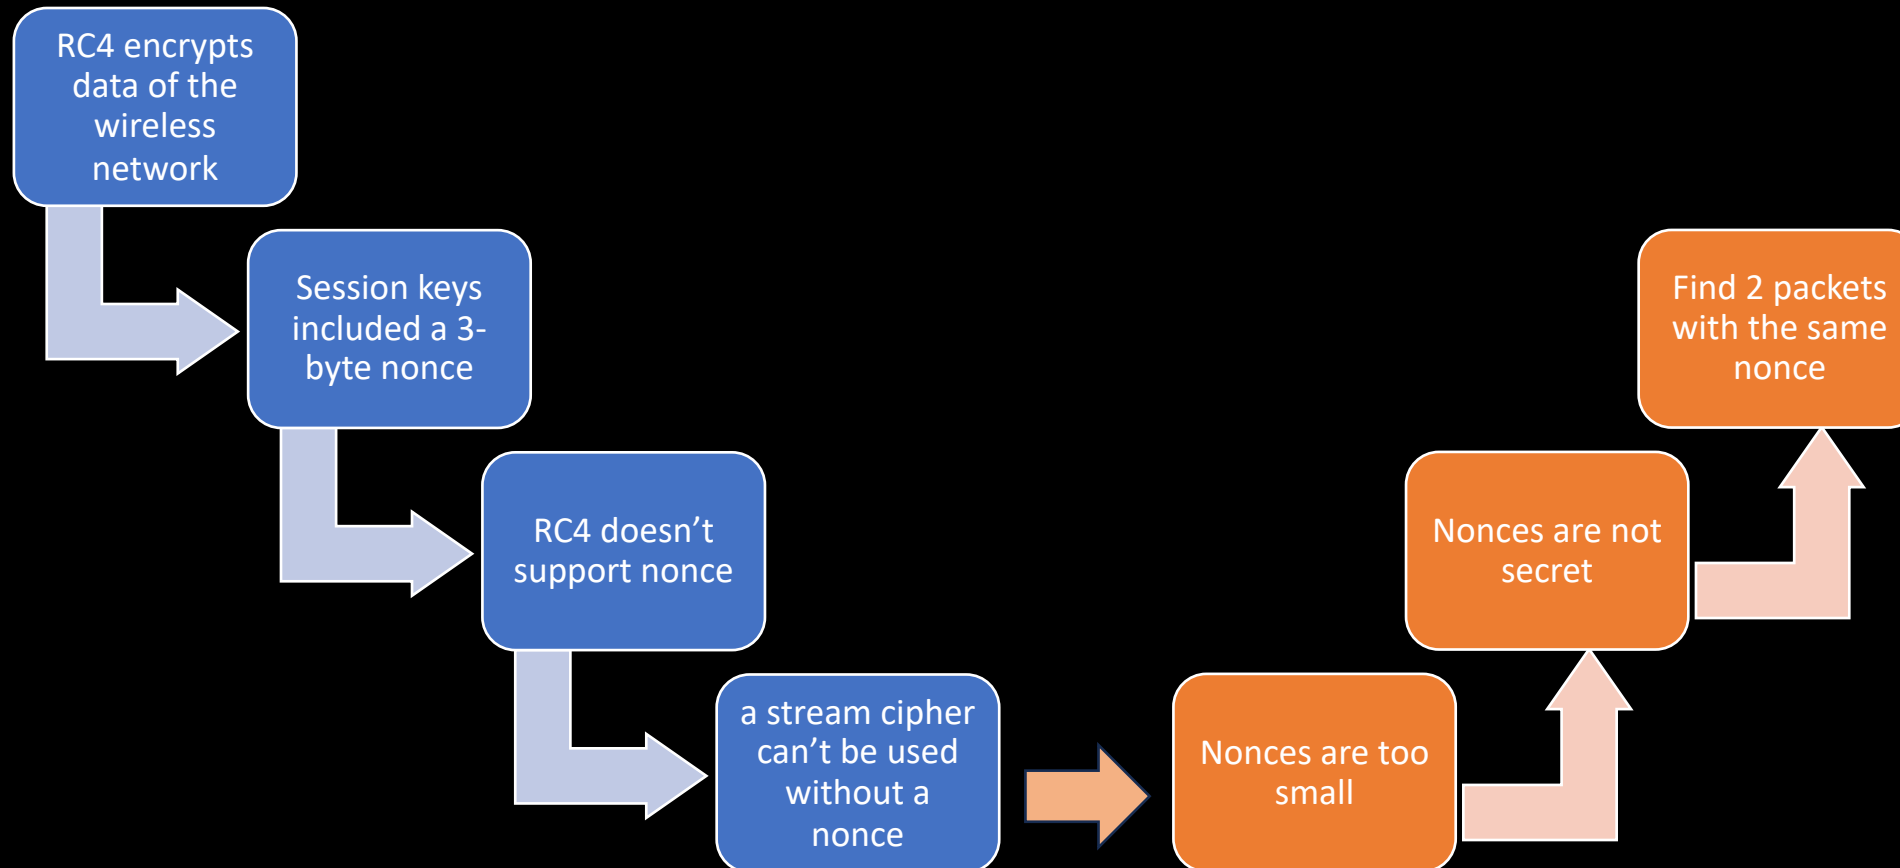
- RC4 in WEP



# Software-Oriented Stream Ciphers

## RC4

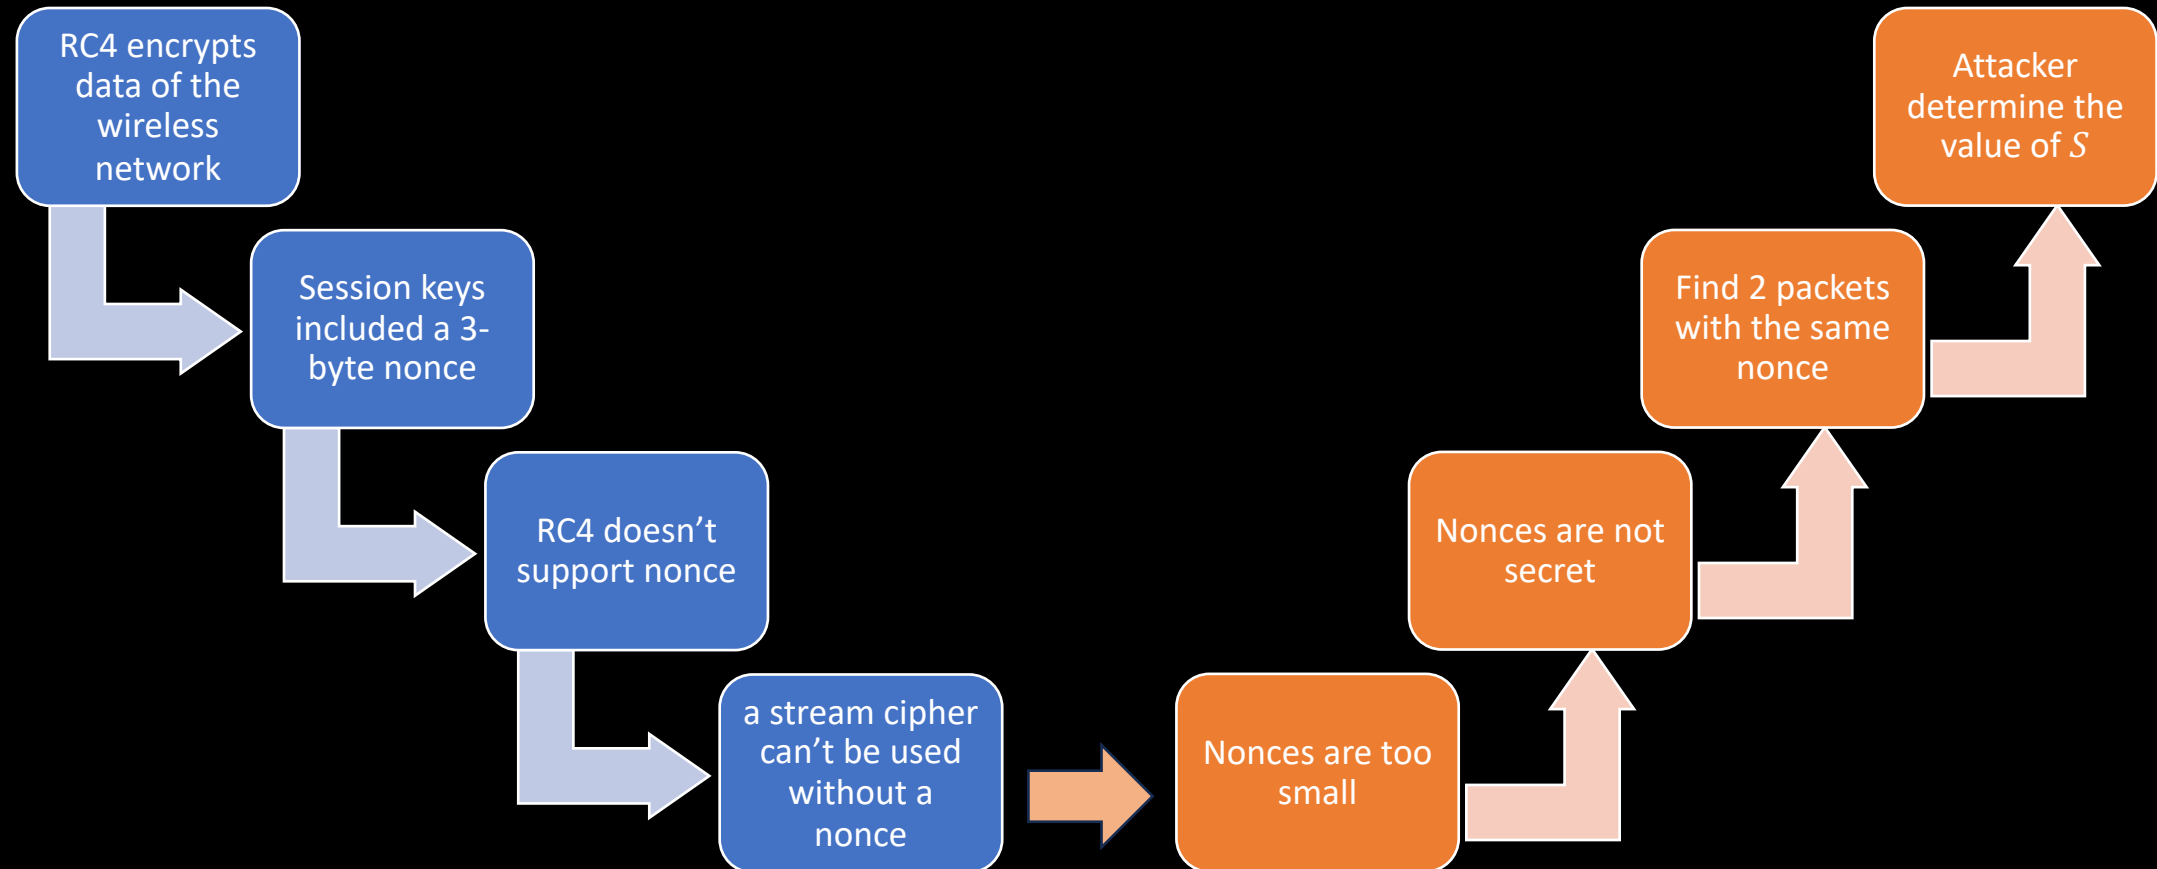
- RC4 in WEP



# Software-Oriented Stream Ciphers

## RC4

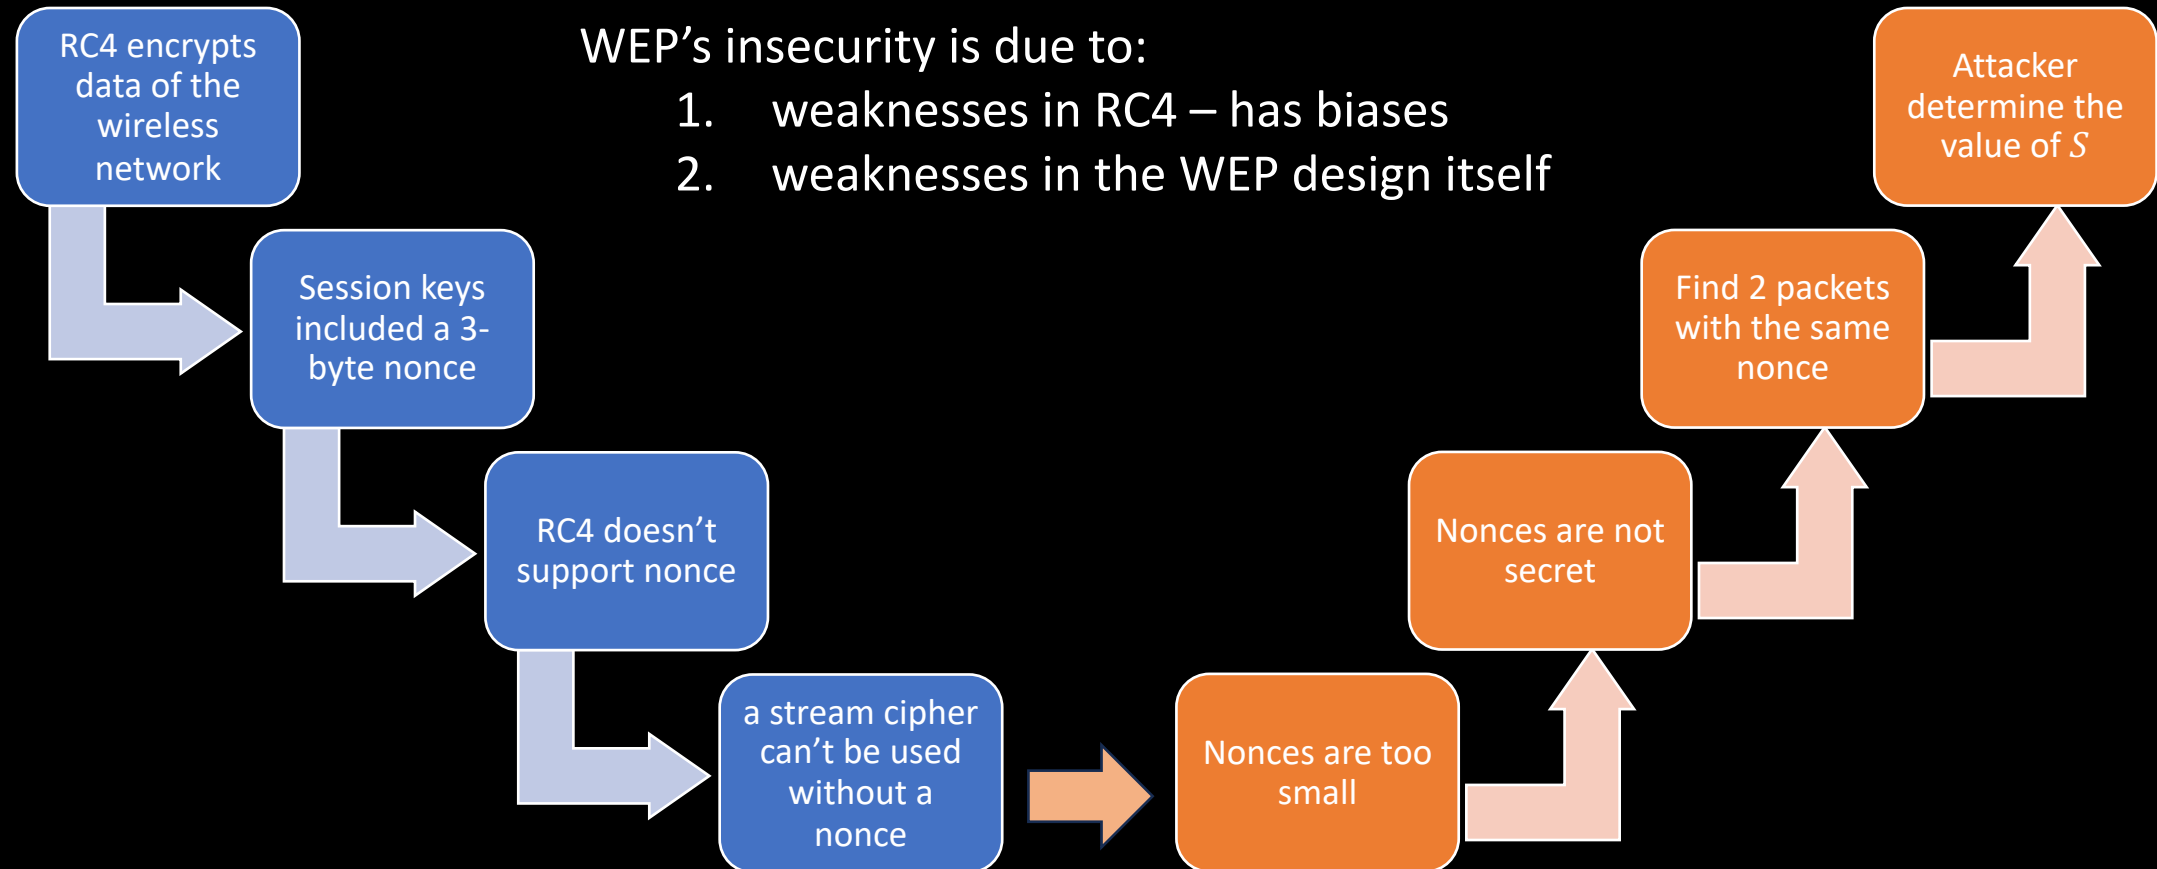
- RC4 in WEP



# Software-Oriented Stream Ciphers

## RC4

- RC4 in WEP



# Software-Oriented Stream Ciphers

## RC4

### RC4 in TLS

- TLS inputs a unique 128-bit session key to the RC4.
- TLS was weak due only to RC4.
  - RC4 has statistical biases
  - Statistical biases are caused by non-randomness
  - Example: the 2<sup>nd</sup> keystream byte is zero, with a probability of  $\frac{1}{128}$ , instead of  $\frac{1}{256}$

# Software-Oriented Stream Ciphers


## RC4

- RC4 in TLS attack: collect many ciphertext and just look for plaintext values
  - The ctexts should be encrypting the same ptxts several times using different keys.
  - This a *broadcast attack model*
- This attack works because of RC4 is biased → the KS is likely to be zeros.
  - Therefore, some ciphertext bytes will be equal to the plaintext bytes.
- Hard to put the attack in practice
  - You need to trick the server to encrypt the same message several times.

# Software-Oriented Stream Ciphers

## Salsa20

- Salsa20 is a counter-based stream cipher
  - It generates its keystream by processing a counter incremented for each block.
- Inputs: 512-bit block of key || nonce || counter
- Transforms the inputs using Salsa20 core function
- The original value of the block is added to the output of the Salsa20 core function to produce the keystream
- The keystream is XORed with the plaintext to get the ciphertext



# Software-Oriented Stream Ciphers

## Salsa20

- Salsa20 core uses a *quarter round (QR)* permutation function.
- The QR transforms four 32-bit words  $a, b, c, d$  as follows:


$$\begin{aligned} b &= b \oplus [(a + d) \lll 7] \\ c &= c \oplus [(b + a) \lll 9] \\ d &= d \oplus [(c + b) \lll 13] \\ a &= a \oplus [(d + c) \lll 18] \end{aligned}$$



# Software-Oriented Stream Ciphers

## Salsa20

- 512-bit block transformation  $\rightarrow$  apply QR to a 4x4 state of 32-bit words.
- The 4x4 internal state consists of:
  - Key: 8 words (256 bits),  $k$
  - Nonce: 2 words (64 bits),  $v$
  - Counter: 2 words (64 bits),  $t$
  - Constants: 4 words (128 bits),  $c$



*Figure 5-11: The initialization of Salsa20's state*

# Software-Oriented Stream Ciphers

## Salsa20

- Salsa20 applies the QR to the columns and the rows of the internal state.
  - **Row round**: transforming all the 4 rows independently
  - **Column round**: transforming all the 4 columns independently
  - The sequence column-round/row-round is called a **double-round**.
- It repeats 10 double-rounds, for 20 rounds in total, thus the 20 in Salsa20.

$\text{QR}(x_0, x_1, x_2, x_3)$   
 $\text{QR}(x_5, x_6, x_7, x_4)$   
 $\text{QR}(x_{10}, x_{11}, x_8, x_9)$   
 $\text{QR}(x_{15}, x_{12}, x_{13}, x_{14})$




Figure 5-12: Columns and rows transformed by Salsa20's quarter-round (QR) function

$\text{QR}(x_0, x_4, x_8, x_{12})$   
 $\text{QR}(x_1, x_5, x_9, x_{13})$   
 $\text{QR}(x_2, x_6, x_{10}, x_{14})$   
 $\text{QR}(x_3, x_7, x_{11}, x_{15})$

# Software-Oriented Stream Ciphers

## Salsa20

- To show why Salsa20 is more secure than RC4, use **differential cryptanalysis**
  - The study of the differences between states rather than their actual values.
- Example: given two states of the Salsa20
  - Both states have the same key (all zeros), nonce (all ff), and constants
  - The first state has counter 0 and the second block has counter one

61707865	00000000	00000000	00000000	61707865	00000000	00000000	00000000
00000000	3320646e	ffffffff	ffffffff	00000000	3320646e	ffffffff	ffffffff
<b>00000000</b>	00000000	79622d32	00000000	<b>00000001</b>	00000000	79622d32	00000000
00000000	00000000	00000000	6b206574	00000000	00000000	00000000	6b206574

*Listing 5-3: Salsa20's initial states for the first two blocks with an all-zero key and an all-one nonce*

# Software-Oriented Stream Ciphers

## Salsa20

Only one bit difference  
between the states

```
61707865 00000000 00000000 00000000
00000000 3320646e ffffffff ffffffff
00000000 00000000 79622d32 00000000
00000000 00000000 00000000 6b206574
```

```
61707865 00000000 00000000 00000000
00000000 3320646e ffffffff ffffffff
00000001 00000000 79622d32 00000000
00000000 00000000 00000000 6b206574
```

After 10 double rounds, everything looks random and no statistical bias  
*In RC4, everything look random, but **had statistical biases***

```
e98680bc f730ba7a 38663ce0 5f376d93
85683b75 a56ca873 26501592 64144b6d
6dcb46fd 58178f93 8cf54cfe cfdc27d7
68bbe09e 17b403a1 38aa1f27 54323fe0
```

```
1ba4d492 c14270c3 9fb05306 ff808c64
b49a4100 f5d8fbbd 614234a0 e20663d1
12e1e116 6a61bc8f 86f01bcb 2efead4a
77775a13 d17b99d5 eb773f5b 2c3a5e7d
```

*Listing 5-4: The states from Listing 5-3 after 10 Salsa20 double-rounds*

# Software-Oriented Stream Ciphers

## Salsa20

### How this works?

If you compute the difference (XOR) between these two states,

61707865	00000000	00000000	00000000	61707865	00000000	00000000	00000000
00000000	3320646e	ffffffff	ffffffff	00000000	3320646e	ffffffff	ffffffff
<b>00000000</b>	00000000	79622d32	00000000	<b>00000001</b>	00000000	79622d32	00000000
00000000	00000000	00000000	6b206574	00000000	00000000	00000000	6b206574

You will get this state

00000000	00000000	00000000	00000000
00000000	00000000	00000000	00000000
<b>00000001</b>	00000000	00000000	00000000
00000000	00000000	00000000	00000000

# Software-Oriented Stream Ciphers

## Salsa20

### How this works?

```
00000000 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00000000 00000000 00000000 00000000
```

# Software-Oriented Stream Ciphers

## Salsa20

### How this works?

```
00000000 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00000000 00000000 00000000 00000000
```

After 1 round



```
80040003 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00002000 00000000 00000000 00000000
```

# Software-Oriented Stream Ciphers

## Salsa20


### How this works?

```
00000000 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00000000 00000000 00000000 00000000
```

After 1 round



```
80040003 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00002000 00000000 00000000 00000000
```



After 1  
round

```
9ed7eb7f 060002c0 18028b0c 57ca83c0
00000000 00000000 00000000 00000000
00000001 0000e000 801c0006 00000000
00002000 00400000 04000008 0060f300
```



# Software-Oriented Stream Ciphers

## Salsa20

### How this works?

```
00000000 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00000000 00000000 00000000 00000000
```

After 1 round




```
80040003 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00002000 00000000 00000000 00000000
```

```
3ab3c25d 9f40a5c9 10070e30 07bd03c0
db1ee2ce 43ee9401 21a702c3 48fd800c
403c1e72 00034003 4dc843be 700b8857
5625b75b 09c00e00 06000348 23f712d4
```

After 1  
round



```
9ed7eb7f 060002c0 18028b0c 57ca83c0
00000000 00000000 00000000 00000000
00000001 0000e000 801c0006 00000000
00002000 00400000 04000008 0060f300
```



After 1  
round

# Software-Oriented Stream Ciphers

## Salsa20

### How this works?

```
00000000 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00000000 00000000 00000000 00000000
```

After 1 round



```
80040003 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00002000 00000000 00000000 00000000
```

After 1  
round



```
d93bed6d a267bf47 760c2f9f 4a41d54b
0e03d792 7340e010 119e6a00 e90186af
7fa9617e b6aca0d7 4f6e9a4a 564b34fd
98be796d 64908d32 4897f7ca a684a2df
```



After 1  
round

```
3ab3c25d 9f40a5c9 10070e30 07bd03c0
db1ee2ce 43ee9401 21a702c3 48fd800c
403c1e72 00034003 4dc843be 700b8857
5625b75b 09c00e00 06000348 23f712d4
```

After 1  
round



```
9ed7eb7f 060002c0 18028b0c 57ca83c0
00000000 00000000 00000000 00000000
00000001 0000e000 801c0006 00000000
00002000 00400000 04000008 0060f300
```

# Software-Oriented Stream Ciphers

## Salsa20

So, after only 4 rounds, a single difference propagates to the 512-bit state.

```
00000000 00000000 00000000 00000000
00000000 00000000 00000000 00000000
00000001 00000000 00000000 00000000
00000000 00000000 00000000 00000000
```



After 4 rounds


```
d93bed6d a267bf47 760c2f9f 4a41d54b
0e03d792 7340e010 119e6a00 e90186af
7fa9617e b6aca0d7 4f6e9a4a 564b34fd
98be796d 64908d32 4897f7ca a684a2df
```

This is called **full diffusion**

# Software-Oriented Stream Ciphers

## Salsa20

- **Attack on Salsa20/8:** exploit a statistical bias in the core algorithm after 4 rounds
  - Salsa20/8 is a smaller version of Salsa20 that uses 8 rounds instead of 20.
  - Published in *New Features of Latin Dances: Analysis of Salsa, ChaCha, and Rumba*
  - Theoretical attack: its complexity is  $2^{251}$  instead of  $2^{256}$



# Software-Oriented Stream Ciphers

## Salsa20

### TASKS

1. Use *pycryptodome* to encrypt a message with *Salsa20* and *ChaCha20* ciphers.

Nonce length	Description	Max data	If random nonce and same key
8 bytes (default)	The original ChaCha20 designed by Bernstein.	No limitations	Max 200 000 messages
12 bytes	The TLS ChaCha20 as defined in <a href="#">RFC7539</a> .	256 GB	Max 13 billions messages
24 bytes	XChaCha20, still in <a href="#">draft stage</a> .	256 GB	No limitations

# Software-Oriented Stream Ciphers

## Salsa20

### 1. Implement a stateful-stream-cipher as follows:

- The *init* function calls a *secure\_PRP* function and returns a *state*
- The *secure\_PRP* function takes a key and a nonce and returns the encryption of the nonce (the initial state). The PRP function is the AES cipher in *ECB* mode.
- The *update* function takes the *previous\_state* as a parameter and calls the *secure\_PRP* function to encrypt the *previous\_state*. The key is 16-bytes 0s.
- Implement the encryption and decryption. The block size is 16 bytes (128 bits)
- What happens if we set the AES mode of operation to CBC mode?

